

In [77]: 1 `import pandas as pd`*#importing pandas*

In [78]: 1 `data=pd.read_csv("/home/placement/Downloads/fiat500.csv")`*#reading csv file*

In [79]: 1 `data.describe()`*#describe data frame*

Out[79]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [80]: 1 data.tail(10)#display top 10 rows
```

Out[80]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
1528	1529	lounge	51	2861	126000	1	43.841980	10.51531	5500
1529	1530	lounge	51	731	22551	1	38.122070	13.36112	9900
1530	1531	lounge	51	670	29000	1	45.764648	8.99450	10800
1531	1532	sport	73	4505	127000	1	45.528511	9.59323	4750
1532	1533	pop	51	1917	52008	1	45.548000	11.54947	9900
1533	1534	sport	51	3712	115280	1	45.069679	7.70492	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.66687	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.41348	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.68227	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.56827	7900

```
In [81]: 1 data1=data.drop(['ID','lat','lon'],axis=1)#deleting column
```

In [82]: 1 data1

Out[82]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

In [83]: 1 data1=pd.get\_dummies(data1)*#dummies of data*

In [84]: 1 data1

Out[84]:

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

In [85]: 1 y=data1['price']*#creating 2 data frames*  
 2 x=data1.drop('price',axis=1)

In [86]:

1 y

Out[86]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
```

Name: price, Length: 1538, dtype: int64

In [87]:

1 x

Out[87]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...	...	...	...	...	...	...	...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [88]: 1 from sklearn.model_selection import train_test_split#testing and training data
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [89]: 1 x_test.head(5)
```

Out[89]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [90]: 1 x_train.shape
```

Out[90]: (1030, 7)

```
In [91]: 1 y_train.shape
```

Out[91]: (1030,)

```
In [92]: 1 x_train.head()
```

Out[92]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [93]: 1 y_train.head()
```

```
Out[93]: 527    9990
         129    9500
         602    7590
         331    8750
         323    9100
         Name: price, dtype: int64
```

```
In [94]: 1 x_test.head()
```

```
Out[94]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [95]: 1 y_test.head()
```

```
Out[95]: 481    7900
         76    7900
         1502   9400
         669   8500
         1409   9700
         Name: price, dtype: int64
```

```
#linear regression
```

```
In [96]: 1 from sklearn.linear_model import LinearRegression
        2 reg=LinearRegression()
        3 reg.fit(x_train,y_train)
```

Out[96]: LinearRegression()

```
In [97]: 1 ypred=reg.predict(x_test)
```

```
In [98]: 1 ypred
```

```
10017.8490121 , 10590.33289679, 10161.75393066, 4927.49556508,
 7276.18410037, 9678.26477249, 9764.65653403, 5643.53722047,
10062.84554534, 5163.04602382, 8307.60791348, 7441.80993846,
 7868.82460983, 9725.36143983, 8669.20982667, 10447.15719448,
 7124.58453563, 9718.32989102, 8059.66615638, 7430.65975056,
10425.57075395, 10364.18738085, 5433.2724385 , 9102.40298437,
 9629.06913727, 10532.3506032 , 10129.42684118, 9149.48843328,
 6158.13422239, 9721.03634157, 10419.02236947, 8838.50241314,
 8182.78836676, 10012.21373766, 9468.92324529, 9904.31954667,
10475.66003551, 10475.0702782 , 9609.27020577, 8115.22501265,
10439.02404036, 10363.81936482, 8720.0683498 , 8274.3579289 ,
 6889.7195761 , 10191.45963957, 4819.0674709 , 8814.11814085,
 5737.62378403, 10051.06593609, 8840.87520652, 10054.31165256,
 9686.269121 , 10463.56977746, 10133.15815395, 9762.80613855,
 9793.03056946, 6796.69068198, 9599.3262671 , 8488.31539047,
 6705.66818403, 10307.58651641, 10045.18332239, 10120.36242166,
 5836.93199112, 8772.49782933, 9680.77538859, 5719.87463854,
 8398.59735084, 9680.77538859, 4334.81943405, 10015.00600846,
 9850.72458719, 7864.73798641, 10072.71245374, 10552.64805598,
10252.47474008, 6861.80726606, 6484.22640656, 10274.62122622
```

```
In [99]: 1 from sklearn.metrics import r2_score
        2 r2_score(y_test,ypred)
```

Out[99]: 0.8415526986865394



```
In [100]: 1 from sklearn.metrics import mean_squared_error as ns
          2 o=ns(y_test,ypred)
          3 o
```

Out[100]: 581887.727391353

```
In [101]: 1 import math
          2 math.sqrt(o)
```

Out[101]: 762.8156575420782

```
In [102]: 1 ypred
```

```
7868.82460983, 9725.36143983, 8669.20982667, 10447.15719448,
7124.58453563, 9718.32989102, 8059.66615638, 7430.65975056,
10425.57075395, 10364.18738085, 5433.2724385 , 9102.40298437,
9629.06913727, 10532.3506032 , 10129.42684118, 9149.48843328,
6158.13422239, 9721.03634157, 10419.02236947, 8838.50241314,
8182.78836676, 10012.21373766, 9468.92324529, 9904.31954667,
10475.66003551, 10475.0702782 , 9609.27020577, 8115.22501265,
10439.02404036, 10363.81936482, 8720.0683498 , 8274.3579289 ,
6889.7195761 , 10191.45963957, 4819.0674709 , 8814.11814085,
5737.62378403, 10051.06593609, 8840.87520652, 10054.31165256,
9686.269121 , 10463.56977746, 10133.15815395, 9762.80613855,
9793.03056946, 6796.69068198, 9599.3262671 , 8488.31539047,
6705.66818403, 10307.58651641, 10045.18332239, 10120.36242166,
5836.93199112, 8772.49782933, 9680.77538859, 5719.87463854,
8398.59735084, 9680.77538859, 4334.81943405, 10015.00600846,
9850.72458719, 7864.73798641, 10072.71245374, 10552.64805598,
10253.47474908, 6861.80736606, 6484.22649656, 10374.62123623,
8426.37409382, 5447.47569851, 9914.20077691, 4687.39013431,
7885.32100747, 5431.00822998, 9911.86294348, 10390.16991322,
6666.64745001, 8844.57015500, 7764.88471004, 4257.54640050
```

```
In [103]: 1 Results=pd.DataFrame(columns=['price','predicted'])
          2 Results['price']=y_test
          3 Results['predicted']=ypred
          4 Results=Results.reset_index()
          5 Results['ID']=Results.index
          6 Results.head(15)
          7
```

Out[103]:

	index	price	predicted	ID
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9
10	1088	9890	10434.349636	10
11	576	7990	7732.262557	11
12	965	7380	7698.672401	12
13	1488	6800	6565.952404	13
14	1432	8900	9662.901035	14

```
In [104]: 1 Results['price_diff']=Results.apply(lambda row: row.price - row.predicted,axis=1)
```

In [105]:

1 Results

Out[105]:

	index	price	predicted	ID	price_diff
0	481	7900	5867.650338	0	2032.349662
1	76	7900	7133.701423	1	766.298577
2	1502	9400	9866.357762	2	-466.357762
3	669	8500	9723.288745	3	-1223.288745
4	1409	9700	10039.591012	4	-339.591012
...	...	...	...	...	...
503	291	10900	10032.665135	503	867.334865
504	596	5699	6281.536277	504	-582.536277
505	1489	9500	9986.327508	505	-486.327508
506	1436	6990	8381.517020	506	-1391.517020
507	575	10900	10371.142553	507	528.857447

508 rows × 5 columns

#ridge regression

```
In [106]: 1 from sklearn.model_selection import GridSearchCV
2 from sklearn.linear_model import Ridge
3 #ridge regression
4 alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]
5
6 ridge = Ridge()
7
8 parameters = {'alpha': alpha}
9
10 ridge_regressor = GridSearchCV(ridge, parameters)
11
12 ridge_regressor.fit(x_train, y_train)
13
```

```
Out[106]: GridSearchCV(estimator=Ridge(),
                        param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                                5, 10, 20, 30]})
```

```
In [107]: 1 ridge_regressor.best_params_
```

```
Out[107]: {'alpha': 30}
```

```
In [108]: 1 ridge=Ridge(alpha=30)
2 ridge.fit(x_train,y_train)
3 y_pred_ridge=ridge.predict(x_test)
```

```
In [109]: 1 from sklearn.metrics import mean_squared_error
2 Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
3 Ridge_Error
```

```
Out[109]: 579521.7970897449
```

```
In [110]: 1 from sklearn.metrics import r2_score
2 r2_score(y_test,y_pred_ridge)
```

```
Out[110]: 0.8421969385523054
```


```
In [111]: 1 Results=pd.DataFrame(columns=['Actual','predicted'])
          2 Results['Actual']=y_test
          3 Results['predicted']=y_pred_ridge
          4 Results=Results.reset_index()
          5 Results['ID']=Results.index
          6 Results.head(10)
          7
```

Out[111]:

	index	Actual	predicted	ID
0	481	7900	5869.741155	0
1	76	7900	7149.563327	1
2	1502	9400	9862.785355	2
3	669	8500	9719.283532	3
4	1409	9700	10035.895686	4
5	1414	9900	9650.311090	5
6	1089	9900	9669.183317	6
7	1507	9950	10115.128380	7
8	970	10700	9900.241944	8
9	1198	8999	9347.080772	9

```
In [112]: 1 import seaborn as sns
          2 import matplotlib.pyplot as plt
          3 sns.lineplot(x='ID',y='Actual',data=Results.head(50))
          4 sns.lineplot(x='ID',y='predicted',data=Results.head(50))
          5 plt.plot()
```

Out[112]: []



#elastic

```
In [113]: 1 from sklearn.linear_model import ElasticNet
          2 from sklearn.model_selection import GridSearchCV
          3 elastic = ElasticNet()
          4
          5 parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
          6
          7 elastic_regressor = GridSearchCV(elastic, parameters)
          8
          9 elastic_regressor.fit(x_train, y_train)
         10
```

```
Out[113]: GridSearchCV(estimator=ElasticNet(),
                        param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                                5, 10, 20]})
```

```
In [114]: 1 elastic_regressor.best_params_
```

```
Out[114]: {'alpha': 0.01}
```

```
In [115]: 1 elastic=ElasticNet(alpha=0.01)
          2 elastic.fit(x_train,y_train)
          3 y_pred_elastic=elastic.predict(x_test)
```

```
In [116]: 1 from sklearn.metrics import r2_score
          2 r2_score(y_test,y_pred_elastic)
```

```
Out[116]: 0.841688021120299
```

```
In [117]: 1 from sklearn.metrics import mean_squared_error
          2 Elasticnet_Error=mean_squared_error(y_pred_elastic,y_test)
          3 Elasticnet_Error
```

```
Out[117]: 581390.7642825295
```

```
In [118]: 1 Results=pd.DataFrame(columns=['Actual','predicted'])
          2 Results['Actual']=y_test
          3 Results['predicted']=y_pred_elastic
          4 Results=Results.reset_index()
          5 Results['ID']=Results.index
          6 Results.head(10)
          7
```

Out[118]:

	index	Actual	predicted	ID
0	481	7900	5867.742075	0
1	76	7900	7136.527402	1
2	1502	9400	9865.726723	2
3	669	8500	9722.573593	3
4	1409	9700	10038.936496	4
5	1414	9900	9653.407122	5
6	1089	9900	9672.438692	6
7	1507	9950	10118.075470	7
8	970	10700	9903.219809	8
9	1198	8999	9350.750929	9



```
In [119]: 1 import seaborn as sns#plot
          2 import matplotlib.pyplot as plt
          3 sns.lineplot(x='ID',y='Actual',data=Results.head(50))
          4 sns.lineplot(x='ID',y='predicted',data=Results.head(50))
          5 plt.plot()
```

Out[119]: []

